id	uuid	product_id	quantity	revenue	balance_earnings	product_data	payout_id	stripe_id
1	9b058043-55fe-4bc7-be5c-50baaf583a98	1	1	3900	3120	{"id": 1, "url": "...", "uuid": "...", "title": "..."}	NULL	acac83e2-5a44-3e27-a588-06f86ce202e6
2	9b058046-37e6-48e1-997c-aeaae908e96d	2	1	3900	3120	...	NULL	...

A few things to consider:

- There's a between `revenue` and `balance_earnings`. Paddle takes 5% from every transaction and your beloved country takes another 5%-27% in the name of VAT. By the way, Hungary (my country) has the highest VAT in the world with 27%. So there are 4 additional columns in the table:
  - `fee_rate` and `fee_amount`. The first one is a percentage and the second one is a dollar amount.
  - `tax_rate` and `tax_amount`.
  - `balance_earnings = revenue - tax_amount - fee_amount`
- `product_data` is a JSON column that stores the whole `Product` object at the moment of the sale. Of course, we store the `product_id` as a foreign key, but we need to store the original product because products can be changed. If product 1 had a price of $39 at the time of the transaction but now it costs $59 it can cause some trouble later. By storing everything in the `product_data` column we always have the source of truth. However, we still need the `product_id` to write optimal queries, such as exporting all transactions by a given product. When you have millions of rows, you probably don't want to use an unindexed JSON column in a `join` statement.
- When a transaction happens we're going to invoke Stripe API and we'll get an ID in response. This is the `stripe_id` column.
- We're going to talk more about `payout_id` later.
- The table also has the following columns: `customer_email`, `user_id`, `created_at`