Summary

The handling of text is essentially the core of what we do as web developers. I hope that this chapter has made it abundantly clear that strings are not magically meaningful to the machines that are processing them; what is ultimately rendered in the end-user's console or web browser is the culmination of layers of rules and your own understanding. Making the mental leap from byte streams to human language is an essential skill for any software engineer that takes their job seriously.

The vast majority of systems on the web today store and transmit strings as UTF-8, and indeed the PSR standards demand that PHP source code be stored in that encoding. Nevertheless, other layers might well opt for UTF-16 (hello JavaScript and older versions of Windows), and it's not unusual for older systems to opt for one of the extended ASCII formats. Hopefully there has been enough material and guidance in this chapter to help you navigate the pitfalls of legacy code. I wish you happy text-crunching!